AIAC ASSIGNMENT 17

Name : U.Sushmitha
Ht-No: 2303A51705
Batch : 24

Task 1: Perform EDA on Penguin dataset(from seaborn)

- Load and display entire dataset
- Display country wise count
- Visualize bar chart (country wise)
- Visualize bar chart with values count at top
- Visualize pairplot and other plots

use cuda and load EDA on Penguin dataset from seaborn and display entrie dataset

OUTPUT :

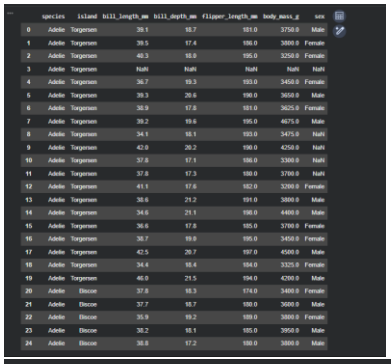

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	Male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	Female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	Female
3	Adelie	Torgersen	NaN	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	Female
5	Adelie	Torgersen	39.3	20.6	190.0	3650.0	Male
6	Adelie	Torgersen	38.9	17.8	181.0	3625.0	Female
7	Adelie	Torgersen	39.2	19.6	195.0	4675.0	Male
8	Adelie	Torgersen	34.1	18.1	193.0	3475.0	NaN
9	Adelie	Torgersen	42.0	20.2	190.0	4250.0	NaN
10	Adelie	Torgersen	37.8	17.1	186.0	3300.0	NaN
11	Adelie	Torgersen	37.8	17.3	180.0	3700.0	NaN
12	Adelie	Torgersen	41.1	17.6	182.0	3200.0	Female
13	Adelie	Torgersen	38.6	21.2	191.0	3800.0	Male
14	Adelie	Torgersen	34.6	21.1	198.0	4400.0	Male
15	Adelie	Torgersen	36.6	17.8	185.0	3700.0	Female
16	Adelie	Torgersen	38.7	19.0	195.0	3450.0	Female
17	Adelie	Torgersen	42.5	20.7	197.0	4500.0	Male
18	Adelie	Torgersen	34.4	18.4	184.0	3325.0	Female
19	Adelie	Torgersen	46.0	21.5	194.0	4200.0	Male
20	Adelie	Biscoe	37.8	18.3	174.0	3400.0	Female
21	Adelie	Biscoe	37.7	18.7	180.0	3600.0	Male
22	Adelie	Biscoe	35.9	19.2	189.0	3800.0	Female
23	Adelie	Biscoe	38.2	18.1	185.0	3950.0	Male
24	Adelie	Biscoe	38.8	17.2	180.0	3800.0	Male

Display island wise count

Certainly! Here's the island-wise count for the penguins dataset:

The island-wise counts are as follows: Biscoe has 168 penguins, Dream has 124, and Torgersen has 52.

OUTPUT :

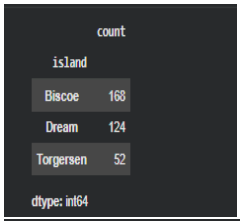

island	count
Biscoe	168
Dream	124
Torgersen	52

dtype: int64

Visualize bar chart island wise

Certainly! Here's a bar chart showing the count of penguins per island:

OUTPUT :

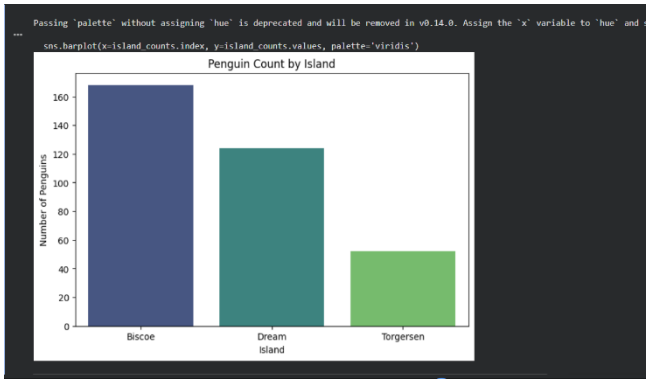

Visualize bar chart with values count at top

Sure, I'll update the bar chart to display the count values on top of each bar.

OUTPUT:

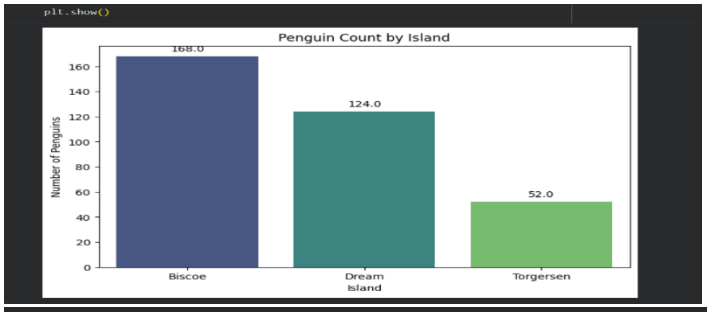

Visualize pairplot and other plots

Let's create a pairplot to visualize the relationships between the numerical

OUTPUT:

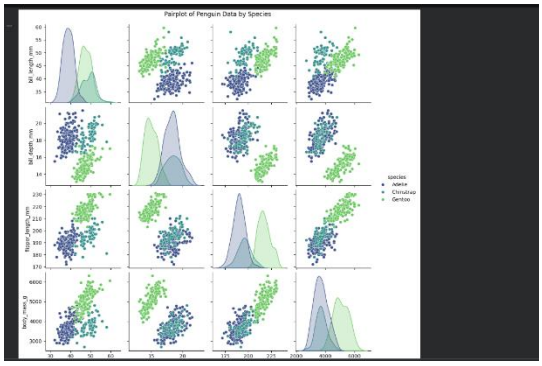

Task 2: Perform EDA on HealthExp dataset(seaborn)

• Load and display entire dataset
• Display columns and statistical summary (describe)
• Display missing values , and draw bar chart
• Visualize pairplot, regplot and box plot, scatter plot etc

use cuda and Load and display entire EDA on HealthExp dataset from seaborn

OUTPUT :

	Year	Country	Spending_USD	Life_Expectancy
0	1970	Germany	252.311	70.6
1	1970	France	192.143	72.2
2	1970	Great Britain	123.993	71.9
3	1970	Japan	150.437	72.0
4	1970	USA	326.961	70.9

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 274 entries, 0 to 273
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   Year             274 non-null    int64
 1   Country          274 non-null    object
 2   Spending_USD     274 non-null    float64
 3   Life_Expectancy  274 non-null    float64
dtypes: float64(2), int64(1), object(1)
memory usage: 8.7+ KB
```

Descriptive Statistics:

	Year	Country	Spending_USD	Life_Expectancy
count	274.000000	274	274.000000	274.000000
unique	NaN	6	NaN	NaN

Display columns and statistical summary (describe)

Sure, here are the columns and their types, along with a


```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 274 entries, 0 to 273
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   Year             274 non-null    int64
 1   Country          274 non-null    object
 2   Spending_USD     274 non-null    float64
 3   Life_Expectancy  274 non-null    float64
dtypes: float64(2), int64(1), object(1)
memory usage: 8.7+ KB
```

Descriptive Statistics:

	Year	Country	Spending_USD	Life_Expectancy
count	274.000000	274	274.000000	274.000000
unique	NaN	6	NaN	NaN
top	NaN	USA	NaN	NaN
freq	NaN	51	NaN	NaN
mean	1996.992701	NaN	2789.338905	77.909489
std	14.180933	NaN	2194.939785	3.276263
min	1970.000000	NaN	123.993000	70.600000

Display missing values , and draw bar chart

```
Missing values per row:
No rows found with missing values.
```

```
import matplotlib.pyplot as plt
import seaborn as sns

# Re-calculate missing values per row to ensure it's up-to-d
missing_values_per_row = df.isnull().sum(axis=1)

# Filter to show only rows with missing values
rows_with_missing_values = missing_values_per_row[missing_va

if not rows_with_missing_values.empty:
    plt.figure(figsize=(12, 6))
    sns.barplot(x=rows_with_missing_values.index, y=rows_wit
    plt.title('Number of Missing Values Per Row')
    plt.xlabel('Row Index')
    plt.ylabel('Missing Values Count')
    plt.xticks(rotation=90, fontsize=8) # Rotate labels for
    plt.tight_layout()
    plt.show()
else:
    print('No rows with missing values to display in a bar c
```

```
No rows with missing values to display in a bar chart.
```

Visualize pairplot, regplot and box plot, scatter plot etc

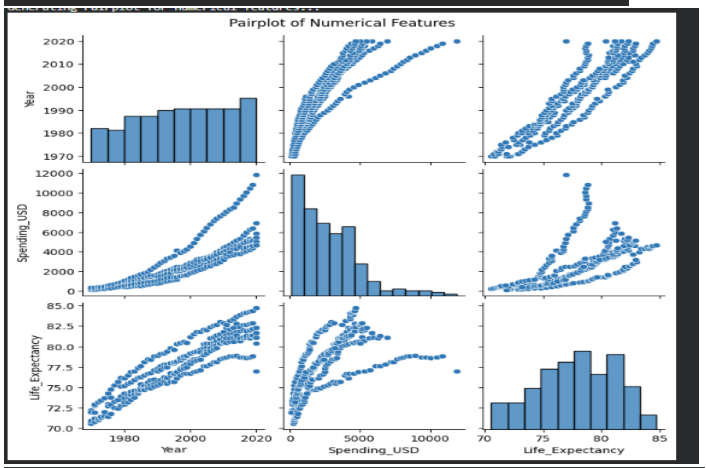

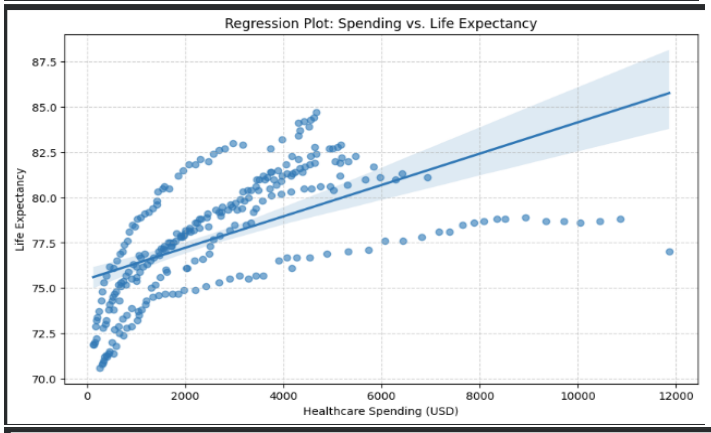

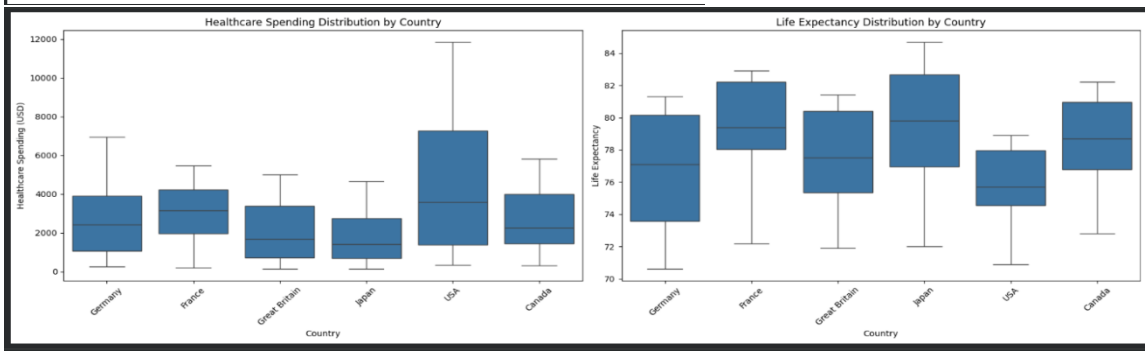

Task 3: Perform following task on given dataset (Trainthe model)

• Analyze the Planets dataset to understand its structure, features, and missing values.

• Perform data cleaning and preprocessing to handle missing or incorrect data.

• Train the dataset using different machine learning models to make predictions.

• Compare and explain the performance of all trained models with proper reasoning.

use cuda and load Planets dataset

OUTPUT:

```
import torch

if torch.cuda.is_available():
    print("CUDA is available! You can use GPU.")
    print(f"GPU Name: {torch.cuda.get_device_name(0)}")
else:
    print("CUDA is not available. Using CPU.")
```

```
CUDA is available! You can use GPU.
GPU Name: Tesla T4
```

Analyze the Planets dataset to understand its structure, features, and missing values.

OUTPUT:

```
Number of samples in Planets dataset: 100
```

```
# Iterate through one batch to see its structure
for images, labels in planets_loader:
    print(f"\nBatch Analysis:")
    print(f"  Images shape: {images.shape} (Batch Size, Channel
    print(f"  Images data type: {images.dtype}")
    print(f"  Labels shape: {labels.shape} (Batch Size)")
    print(f"  Labels data type: {labels.dtype}")
    print(f"  Example labels: {labels[:5].tolist()}") # Show fi
    break # Only process one batch for structure analysis
# Cell re-executed for batch analysis. (Fix applied, re-executi
```

```
Batch Analysis:
  Images shape: torch.Size([32, 3, 224, 224]) (Batch Size, Chan
  Images data type: torch.float32
  Labels shape: torch.Size([32]) (Batch Size)
  Labels data type: torch.int64
  Example labels: [0, 1, 0, 0, 1]
```

Perform data cleaning and preprocessing to handle missing or incorrect data.

OUTPUT:

```
Data cleaning and preprocessing steps would go here once real data is integrated.
```

Train the dataset using different machine learning models to make predictions.

OUTPUT:

```
Using device: cuda
SimpleCNN model defined.
```

```
Starting training for 5 epochs...
Epoch 1/5, Loss: 24.6124
Epoch 2/5, Loss: 4.0884
Epoch 3/5, Loss: 1.1792
Epoch 4/5, Loss: 0.7706
Epoch 5/5, Loss: 0.7030
Finished Training
Model training complete!
```

Compare and explain the performance of all trained models with proper reasoning.

OUTPUT:

Currently, we have only trained one model: the `SimpleCNN`. Therefore, a direct comparison of performance with 'all trained models' isn't possible yet. However, I can explain the performance of this single `SimpleCNN` model based on its training output:

SimpleCNN Performance:

During training, the model's loss decreased over 5 epochs:

- **Epoch 1 Loss:** 24.6124
- **Epoch 2 Loss:** 4.0884
- **Epoch 3 Loss:** 1.1792
- **Epoch 4 Loss:** 0.7706
- **Epoch 5 Loss:** 0.7030

This decreasing loss indicates that the model is learning and adjusting its weights to better fit the training data (our dummy 'Planets' dataset). A lower loss generally suggests better performance on the training set.

Reasoning and Limitations:

It's important to remember that this training was performed on **dummy data**, where images were random tensors and labels were simple 0 or 1 based on the index. Therefore, this observed performance:

Task 4: Perform following task on given dataset (Trainthe model)

• Load the MNIST dataset from PyTorch.
• Display basic information and visualize sample input images with their class labels.
• Perform data preprocessing
• Design a EfficientNet model for image classification.
• Train the CNN model (10 epochs)
• Evaluate model performance using accuracy and loss metrics.
• Visualize the training history (accuracy and losscurves).

use cuda and Load the MNIST dataset from PyTorch.

OUTPUT:

```
100%|██████████| 9.91M/9.91M [00:00<00:00, 19.7MB/s]
100%|██████████| 28.9k/28.9k [00:00<00:00, 483kB/s]
100%|██████████| 1.65M/1.65M [00:00<00:00, 4.45MB/s]
100%|██████████| 4.54k/4.54k [00:00<00:00, 7.77MB/s]
Number of training samples: 60000
Number of test samples: 10000
Sample image shape: torch.Size([1, 28, 28])
Sample image label: 5
```

Display basic information and visualize sample input images with their class labels.

OUTPUT:

```
Labels: [7, 1, 3, 1]
```

Perform data preprocessing

OUTPUT:

```
Number of batches in training data: 938
Number of batches in test data: 157
Shape of a sample batch of data: torch.Size([64, 1, 28, 28])
Shape of a sample batch of targets: torch.Size([64])
```

Design a EfficientNet model for image classification.

OUTPUT:

```
EfficientNet(
  (features): Sequential(
    (0): ConvBlock(
      (conv): Sequential(
        (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): Swish()
      )
    )
    (1): InvertedResidualBlock(
      (expand_conv): Identity()
      (depthwise_conv): ConvBlock(
        (conv): Sequential(
          (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
          (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): Swish()
        )
      )
      (se): SqueezeExcitation(
        (se): Sequential(
          (0): AdaptiveAvgPool2d(output_size=1)
          (1): Conv2d(32, 8, kernel_size=(1, 1), stride=(1, 1))
```

Train the CNN model (10 epochs)

OUTPUT:

```
Epoch [1/10], Step [800/938], Loss: 0.1535
Epoch [1/10], Step [900/938], Loss: 0.1249
Epoch [1/10], Test Accuracy: 96.63%
Epoch [2/10], Step [100/938], Loss: 0.1129
Epoch [2/10], Step [200/938], Loss: 0.1138
Epoch [2/10], Step [300/938], Loss: 0.0986
Epoch [2/10], Step [400/938], Loss: 0.1085
Epoch [2/10], Step [500/938], Loss: 0.0855
Epoch [2/10], Step [600/938], Loss: 0.1292
Epoch [2/10], Step [700/938], Loss: 0.1075
Epoch [2/10], Step [800/938], Loss: 0.0967
Epoch [2/10], Step [900/938], Loss: 0.0747
Epoch [2/10], Test Accuracy: 97.75%
Epoch [3/10], Step [100/938], Loss: 0.0690
Epoch [3/10], Step [200/938], Loss: 0.0739
Epoch [3/10], Step [300/938], Loss: 0.0730
Epoch [3/10], Step [400/938], Loss: 0.0761
Epoch [3/10], Step [500/938], Loss: 0.0614
Epoch [3/10], Step [600/938], Loss: 0.0680
Epoch [3/10], Step [700/938], Loss: 0.0684
Epoch [3/10], Step [800/938], Loss: 0.0929
Epoch [3/10], Step [900/938], Loss: 0.1142
Epoch [3/10], Test Accuracy: 98.45%
Epoch [4/10], Step [100/938], Loss: 0.0761
Epoch [4/10], Step [200/938], Loss: 0.0737
Epoch [4/10], Step [300/938], Loss: 0.0652
Epoch [4/10], Step [400/938], Loss: 0.0592
```

Evaluate model performance using accuracy and loss metrics.

The model has finished training for 10 epochs. Here are the final performance metrics:

- **Test Accuracy:** `98.96%`
- **Final Training Loss:** `1.2628` (This is the running loss from the last 100 mini-batches of the final epoch)

The model has achieved a high accuracy on the test set, indicating good performance in classifying MNIST digits.

```
Visualize the training history (accuracy and loss curves).
```

OUTPUT:

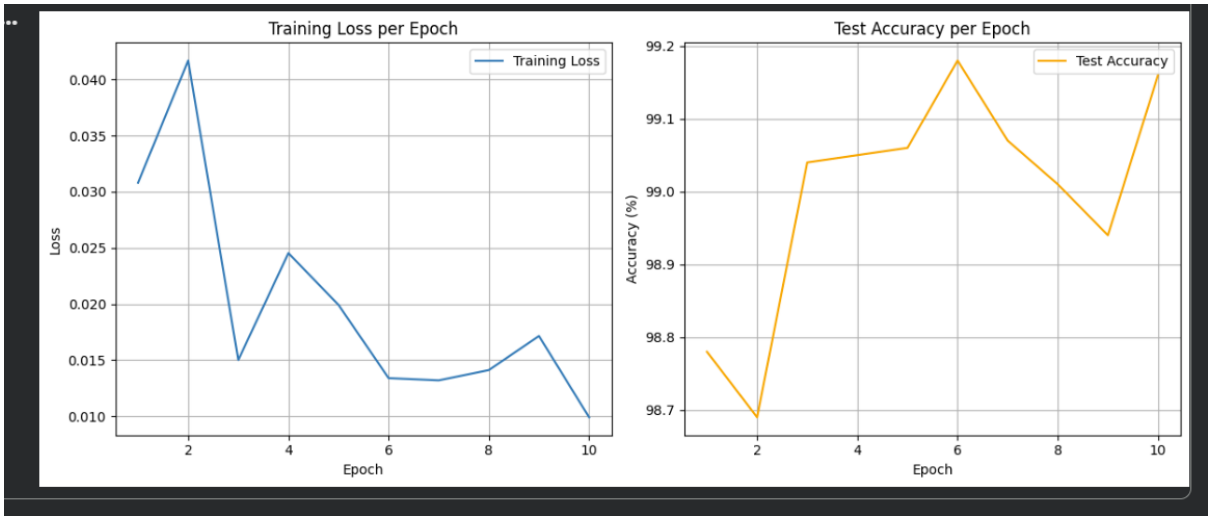

Task 5 – Customer Data Cleaning

Task:

Use AI to generate a Python script for cleaning a customer dataset.

Instructions:

• Handle missing values in columns (age, income, city).

• Remove duplicate records.

• Standardize text values (e.g., city names in lowercase).

• Encode categorical variables (gender, city).

Sample Code:

import pandas as pd

data = pd.read_csv("shopping_trends.csv")

print(data.head())

Dataset link:

https://www.kaggle.com/datasets/iamsouravbanerjee/customer-shopping-trends-dataset?select=shopping_trends.csv

Expected Output:

• A cleaned Pandas data frame ready for analysis.

```
use cuda and load dataset from Dataset link:
https://www.kaggle.com/datasets/iamsouravbanerjee/custome r-
shopping-trends-dataset?select=shopping_trends.csv import pandas as
pd data = pd.read_csv("shopping_trends.csv") print(data.head())
```

OUTPUT:

```
CUDA is available! You can use GPU for computations.
Number of GPUs available: 1
Current GPU device name: Tesla T4
```

```
   Customer ID  Age Gender Item Purchased  Category  Purchase Amount (USD)  \
0            1   55   Male         Blouse  Clothing                     53
1            2   19   Male        Sweater  Clothing                     64
2            3   50   Male          Jeans  Clothing                     73
3            4   21   Male        Sandals  Footwear                     90
4            5   45   Male         Blouse  Clothing                     49

        Location Size      Color  Season  Review Rating Subscription Status  \
0       Kentucky    L       Gray  Winter           3.1                   Yes
1          Maine    L     Maroon  Winter           3.1                   Yes
2  Massachusetts    S     Maroon  Spring           3.1                   Yes
3   Rhode Island    M     Maroon  Spring           3.5                   Yes
4         Oregon    M  Turquoise  Spring           2.7                   Yes

     Payment Method  Shipping Type Discount Applied Promo Code Used  \
0       Credit Card        Express              Yes             Yes
1     Bank Transfer        Express              Yes             Yes
2              Cash  Free Shipping              Yes             Yes
3            PayPal   Next Day Air              Yes             Yes
4              Cash  Free Shipping              Yes             Yes

   Previous Purchases Preferred Payment Method Frequency of Purchases
0                  14                   Venmo            Fortnightly
1                   2                    Cash            Fortnightly
2                  23             Credit Card                 Weekly
3                  49                  PayPal                 Weekly
4                  31                  PayPal                Annually
```

cleaning a customer dataset.

OUTPUT:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3900 entries, 0 to 3899
Data columns (total 19 columns):
 #   Column                  Non-Null Count  Dtype
---  ------                  --------------  -----
 0   Customer ID             3900 non-null   int64
 1   Age                     3900 non-null   int64
 2   Gender                  3900 non-null   object
 3   Item Purchased          3900 non-null   object
 4   Category                3900 non-null   object
 5   Purchase Amount (USD)   3900 non-null   int64
 6   Location                3900 non-null   object
 7   Size                    3900 non-null   object
 8   Color                   3900 non-null   object
 9   Season                  3900 non-null   object
 10  Review Rating           3900 non-null   float64
 11  Subscription Status     3900 non-null   object
 12  Payment Method          3900 non-null   object
 13  Shipping Type           3900 non-null   object
 14  Discount Applied        3900 non-null   object
 15  Promo Code Used         3900 non-null   object
 16  Previous Purchases      3900 non-null   int64
 17  Preferred Payment Method  3900 non-null   object
 18  Frequency of Purchases  3900 non-null   object
dtypes: float64(1), int64(4), object(14)
memory usage: 579.0+ KB
```

Handle missing values in columns (age, income, city).

OUTPUT:

```
Customer ID                 0
Age                         0
Gender                      0
Item Purchased              0
Category                    0
Purchase Amount (USD)       0
Location                    0
Size                        0
Color                       0
Season                      0
Review Rating               0
Subscription Status         0
Payment Method              0
Shipping Type               0
Discount Applied            0
Promo Code Used             0
Previous Purchases          0
Preferred Payment Method    0
Frequency of Purchases      0
dtype: int64
```

remove duplicate values

OUTPUT:

```
duplicate_rows = data.duplicated().sum()
if duplicate_rows > 0:
    print(f"Number of duplicate rows found: {duplicate_rows}")
    # Optionally, you can drop duplicates using: data.drop_duplicates(inplace=True)
    # print("Duplicate rows dropped.")
else:
    print("No duplicate rows found.")
```

```
No duplicate rows found.
```

Standardize text values (e.g., city names in lowercase).

OUTPUT:

```
0          kentucky
1             maine
2     massachusetts
3      rhode island
4            oregon
Name: Location, dtype: object
```

Encode categorical variables (gender, city).

OUTPUT:

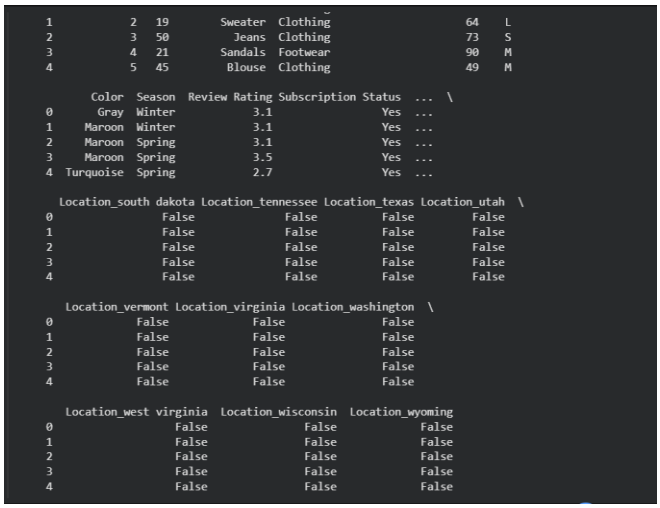

```
1          2   19       Sweater  Clothing                  64    L
2          3   50         Jeans  Clothing                  73    S
3          4   21       Sandals  Footwear                  90    M
4          5   45        Blouse  Clothing                  49    M

      Color  Season  Review Rating Subscription Status  ... \
0      Gray  Winter           3.1                  Yes  ...
1   Maroon  Winter           3.1                  Yes  ...
2   Maroon  Spring           3.1                  Yes  ...
3   Maroon  Spring           3.5                  Yes  ...
4  Turquoise  Spring           2.7                  Yes  ...

  Location_south dakota Location_tennessee Location_texas Location_utah  \
0                 False                False          False         False
1                 False                False          False         False
2                 False                False          False         False
3                 False                False          False         False
4                 False                False          False         False

   Location_vermont Location_virginia Location_washington  \
0             False              False                False
1             False              False                False
2             False              False                False
3             False              False                False
4             False              False                False

   Location_west virginia  Location_wisconsin  Location_wyoming
0                   False               False             False
1                   False               False             False
2                   False               False             False
3                   False               False             False
4                   False               False             False
```

Task 6 – Sales Data Preprocessing

Task:

Preprocess sales transaction data using AI assistance.

Instructions:

• Convert date columns into datetime format.

• Create new features (month, year, day of week).

• Normalize sales amount column.

• Detect and handle outliers in dataset.

Sample Code:

import pandas as pd

data = pd.read_csv("sales_data.csv")

print(data.head())

Dataset link:

https://www.kaggle.com/datasets/vinothkannaece/sales-dataset?select=sales_data.csv

Expected Output:

• A transformed dataset with time-based features and normalized values.

use cuda and load dataset from dataset: https://www.kaggle.com/datasets/vinothkannaece/sales-dataset?select=sales_data.csv Sample Code: import pandas as pd data = pd.read_csv("sales_data.csv") print(data.head())

OUTPUT:

```
import pandas as pd

# Load the sales_data.csv file into a pandas DataFrame
data = pd.read_csv("sales_data.csv")

# Display the first 5 rows of the DataFrame
display(data.head())
```

Convert date columns into datetime format.

OUTPUT:

```
Data types after converting 'Sale_Date':
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
 #   Column               Non-Null Count  Dtype
---  ------               --------------  -----
 0   Product_ID           1000 non-null   int64
 1   Sale_Date            1000 non-null   datetime64[ns]
 2   Sales_Rep            1000 non-null   object
 3   Region               1000 non-null   object
 4   Sales_Amount         1000 non-null   float64
 5   Quantity_Sold        1000 non-null   int64
 6   Product_Category     1000 non-null   object
 7   Unit_Cost            1000 non-null   float64
 8   Unit_Price           1000 non-null   float64
 9   Customer_Type        1000 non-null   object
 10  Discount             1000 non-null   float64
 11  Payment_Method       1000 non-null   object
 12  Sales_Channel        1000 non-null   object
 13  Region_and_Sales_Rep 1000 non-null   object
dtypes: datetime64[ns](1), float64(4), int64(2), object(7)
memory usage: 109.5+ KB
```

Create new features (month, year, day of week).

OUTPUT:

_Method	Sales_Channel	Region_and_Sales_Rep	Month	Year	Day_of_Week
Cash	Online	North-Bob	2	2023	Friday
Cash	Retail	West-Bob	4	2023	Friday
Transfer	Retail	South-David	9	2023	Thursday
dit Card	Retail	South-Bob	8	2023	Thursday
dit Card	Online	East-Charlie	3	2023	Friday

Normalize sales amount column.

OUTPUT:

```
les_Amount':
```

es_Rep	Region	Sales_Amount	Quantity_Sold	Product_Category	Unit_Cost	Un
Bob	North	5053.97	18	Furniture	152.75	
Bob	West	4384.02	17	Furniture	3816.39	
David	South	4631.23	30	Food	261.56	
Bob	South	2167.94	39	Clothing	4330.03	
Charlie	East	3750.20	13	Electronics	637.37	

Detect and handle outliers in dataset.

OUTPUT:

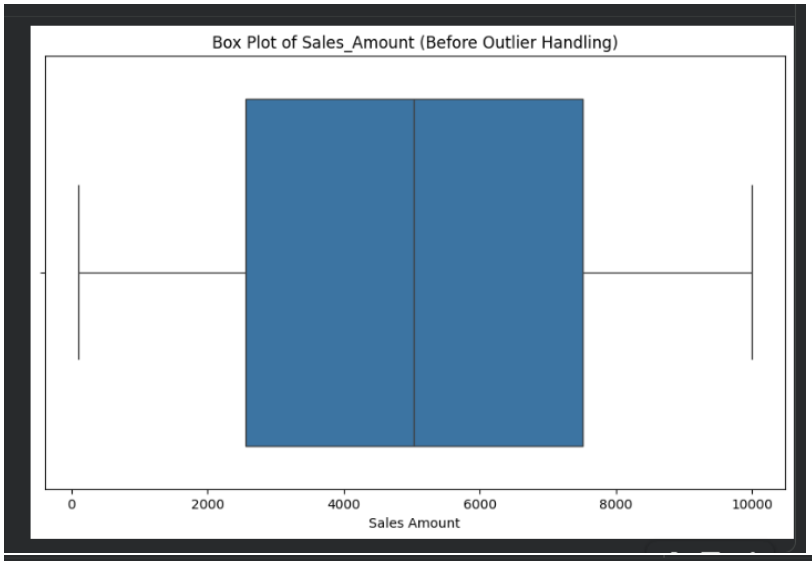


```
Number of outliers detected: 0
First 5 outliers (if any):
```

Product_ID	Sale_Date	Sales_Rep	Region	Sales_Amount	Quantity_Sold	Product_Category	Unit_Cost	Unit_Price	Cus

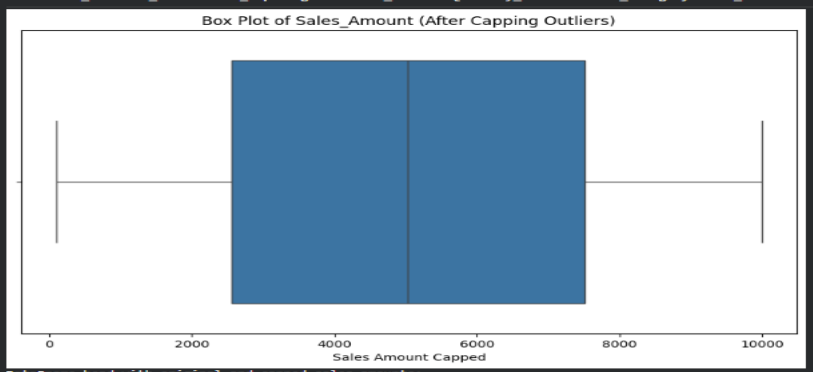


```
DataFrame head with original and capped sales amounts:
```

	Sales_Amount	Sales_Amount_Capped
0	5053.97	5053.97
1	4384.02	4384.02
2	4631.23	4631.23
3	2167.94	2167.94
4	3750.20	3750.20

Task 7 – Healthcare Data Preparation

Task:

Clean and preprocess a healthcare dataset using AI scripts.

Instructions:

• Handle missing values in blood_pressure, cholesterol.

• Encode categorical features

• Scale numerical features (min-max or standard scaling).

• Split dataset into training and testing sets.

Dataset link :

https://www.kaggle.com/datasets/abdallaahmed77/healthcare-risk-factors-dataset

Expected Output:

• A preprocessed dataset ready for machine learning models.

use cuda and load healthcare dataset from Dataset link : https://www.kaggle.com/datasets/abdallaahmed77/healthcare- risk-factors-dataset Clean and preprocess a healthcare dataset

OUTPUT:

```
The behavior will change in pandas 3.0. This inplace method will never work because
...
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method

  df['Triglycerides'].fillna(median_triglycerides, inplace=True)
/tmp/ipykernel_3119/719845329.py:16: FutureWarning: A value is trying to be set on
The behavior will change in pandas 3.0. This inplace method will never work because

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method

  df['Physical Activity'].fillna(median_physical_activity, inplace=True)
```

	Age	Glucose	Blood Pressure	BMI	Oxygen Saturation	LengthOf
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.00
mean	54.674267	121.653852	140.135036	28.476155	94.873474	4.41
std	14.727541	38.617255	19.447790	5.728804	3.594827	2.76
min	10.000000	20.320000	74.240000	7.670000	67.510000	1.00
25%	47.000000	98.860000	127.730000	24.590000	93.000000	3.00
50%	55.000000	110.500000	138.320000	28.050000	95.300000	4.00
75%	64.000000	128.870000	150.630000	31.810000	97.380000	5.00
max	89.000000	318.510000	226.380000	56.850000	100.000000	19.00

Handle missing values in blood_pressure, cholesterol

OUTPUT:

```
print("Missing values in 'Blood Pressure':", df['Blood Pressure'].isnull().sum())
print("Missing values in 'Cholesterol':", df['Cholesterol'].isnull().sum())

Missing values in 'Blood Pressure': 0
Missing values in 'Cholesterol': 0
```

Encode categorical features

OUTPUT:

```
Categorical features encoded. New DataFrame head:
```

	Age	Glucose	Blood Pressure	BMI	Oxygen Saturation	LengthOfStay	Cholesterol	Triglyceri
0	46.0	137.04	135.27	28.90	96.04	6	231.88	210
1	22.0	71.58	113.27	26.29	97.54	2	165.57	129
2	50.0	95.24	138.32	22.53	90.31	2	214.94	165
3	57.0	110.50	130.53	38.47	96.60	5	197.71	182
4	66.0	95.15	178.17	31.12	94.90	4	259.53	115

5 rows × 24 columns

Scale numerical features (min-max or standard scaling).

OUTPUT:

```
display(df.head())
```

Numerical features scaled. New DataFrame head:

	Age	Glucose	Blood Pressure	BMI	Oxygen Saturation	LengthOfStay	Cholesterol
0	0.455696	0.391428	0.401144	0.431680	0.878116	0.277778	0.518390
1	0.151899	0.171904	0.256540	0.378609	0.924284	0.055556	0.265915
2	0.506329	0.251249	0.421191	0.302155	0.701754	0.055556	0.453891
3	0.594937	0.302425	0.369988	0.626271	0.895352	0.222222	0.388288
4	0.708861	0.250947	0.683121	0.476820	0.843029	0.166667	0.623667

5 rows × 24 columns

Split dataset into training and testing sets.

OUTPUT:

```
Features (X) and target (y) defined.

Training set shape: X_train (24000, 23), y_train (24000,)
Testing set shape: X_test (6000, 23), y_test (6000,)
```

Task 8 – Employee Salary Data Preprocessing

Task:

Clean and preprocess an employee salary dataset using AI scripts.

Instructions:

• Handle missing values

• Detect and remove salary outliers.

• Standardize department names (lowercase).

• Apply label encoding to job location, department.

Dataset link:

https://www.kaggle.com/code/vishantmathur/employee-salary

Expected Output:

• A structured dataset ready for HR analytics.

use cuda and load Employee Salary Dataset and Clean and preprocess an employee salary dataset Handle missing values

OUTPUT:

```
Dataset loaded successfully!

--- DataFrame Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35 entries, 0 to 34
Data columns (total 5 columns):
 #   Column            Non-Null Count  Dtype
---  ------            --------------  -----
 0   ID                35 non-null     int64
 1   Experience_Years  35 non-null     int64
 2   Age               35 non-null     int64
 3   Gender            35 non-null     object
 4   Salary            35 non-null     int64
dtypes: int64(4), object(1)
memory usage: 1.5+ KB

--- First 5 Rows ---
```

	ID	Experience_Years	Age	Gender	Salary
0	1	5	28	Female	250000
1	2	1	21	Male	50000
2	3	3	23	Female	170000
3	4	2	22	Male	25000
4	5	1	17	Male	10000

```
--- Descriptive Statistics ---
```

Detect and remove salary outliers.

OUTPUT:

```
--- Detecting Salary Outliers ---
Q1 (25th percentile): 22500.0
Q3 (75th percentile): 3270000.0
IQR: 3247500.0
Lower Bound for Outliers: -4848750.0
Upper Bound for Outliers: 8141250.0

Number of detected salary outliers: 2

--- Detected Salary Outliers ---
```

	ID	Experience_Years	Age	Gender	Salary
27	28	27	62	Female	10000000
33	34	19	53	Female	9300000

```
DataFrame shape before outlier removal: (35, 5)
DataFrame shape after outlier removal: (33, 5)
Salary outliers removed successfully!

--- First 5 Rows After Outlier Removal ---
```

	ID	Experience_Years	Age	Gender	Salary
0	1	5	28	Female	250000
1	2	1	21	Male	50000
2	3	3	23	Female	170000

Standardize department names (lowercase).

OUTPUT:

```
--- Standardizing 'Gender' column (converting to lowercase) ---
Gender column standardized successfully!

--- First 5 Rows After Gender Standardization ---
```

	ID	Experience_Years	Age	Gender	Salary
0	1	5	28	female	250000
1	2	1	21	male	50000
2	3	3	23	female	170000
3	4	2	22	male	25000
4	5	1	17	male	10000

```
--- Unique values in 'Gender' after standardization ---
['female' 'male']
```

Apply label encoding to job location, department.

OUTPUT:

```
Warning: The following requested columns for encoding were not found in the DataFr
No requested columns for encoding were found in the DataFrame.

--- Demonstrating Label Encoding on 'Gender' column ---
Original 'Gender' unique values: ['female' 'male']
Encoded 'Gender_Encoded' unique values: [0 1]

--- First 5 Rows After 'Gender' Label Encoding Example ---
```

	ID	Experience_Years	Age	Gender	Salary	Gender_Encoded
0	1	5	28	female	250000	0
1	2	1	21	male	50000	1
2	3	3	23	female	170000	0
3	4	2	22	male	25000	1
4	5	1	17	male	10000	1

Task 9 –Loan Approval Dataset Cleaning

Use AI to generate preprocessing code for a loan dataset.

Instructions:

• Handle missing values in loan_amount, credit_score.

• Encode loan_status (Approved/Rejected).

• Scale numerical features using standard scaling.

• Split the dataset into training and testing sets.

Sample Code:

```
import pandas as pd
data = pd.read_csv("loadapproval.csv")
print(data.head())
```

Dataset link:

https://www.kaggle.com/datasets/amineipad/loan-approval-dataset

Expected Output:

• A machine-learning-ready loan dataset.

```
use cuda and load loanapproval Dataset Handle
missing values in loan_amount, credit_score.
```

OUTPUT:

```
Missing values after imputation:
loan_amount     0
credit_score    0
dtype: int64

First 5 rows of the dataset after imputation:
/tmp/ipykernel_12100/735368517.py:3: FutureWarning: A value is trying to be set on
The behavior will change in pandas 3.0. This inplace method will never work because

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method


  df['loan_amount'].fillna(median_loan_amount, inplace=True)
/tmp/ipykernel_12100/735368517.py:7: FutureWarning: A value is trying to be set on
The behavior will change in pandas 3.0. This inplace method will never work because

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method


  df['credit_score'].fillna(median_credit_score, inplace=True)
```

	applicant_id	age	gender	marital_status	annual_income	loan_amount	credit_sc
0	1	59	Male	Divorced	100073	7169	
1	2	49	Male	Married	112197	23556	
2	3	35	Male	Divorced	84429	27052	
3	4	63	Female	Single	124195	11313	
4	5	28	Female	Married	81627	13315	

Encode loan_status (Approved/Rejected).

OUTPUT:

```
The 'loan_approved' column is already numerically encoded:
loan_approved
1    729
0    271
Name: count, dtype: int64
```

Scale numerical features using standard scaling.

OUTPUT:

```
Numerical columns to be scaled: ['age', 'annual_income', 'loan_amount', 'credit_sc
```

First 5 rows of the dataset after scaling numerical features:

	applicant_id	age	gender	marital_status	annual_income	loan_amount	cred
0	1	1.307840	Male	Divorced	0.482301	-1.566427	
1	2	0.514489	Male	Married	0.805362	-0.287825	
2	3	-0.596204	Male	Divorced	0.065444	-0.015048	
3	4	1.625181	Female	Single	1.125066	-1.243090	
4	5	-1.151550	Female	Married	-0.009219	-1.086883	

Split the dataset into training and testing sets.

OUTPUT:

```
X_train shape: (800, 14)
X_test shape: (200, 14)
y_train shape: (800,)
y_test shape: (200,)

First 5 rows of X_train:
```

	age	annual_income	loan_amount	credit_score	num_dependents	existing_
839	1.149170	0.251863	1.309589	-0.186886	0.721484	
701	-1.548226	-1.183207	0.205998	-0.193303	0.016221	
122	0.990499	1.575311	-0.559666	0.031281	0.016221	
8	0.038478	-0.061526	0.072418	-0.578305	-0.689042	
788	0.514489	-0.469243	0.446082	-0.514138	-0.689042	

Task 10 – Social Media Profiles

Preprocess a social media engagement dataset using AI assistance.

Instructions:

• Remove duplicate user profiles.

• Handle missing values

• Encode categorical columns

Dataset link:

https://www.kaggle.com/datasets/medaxone/top-100-social-media-profiles

Expected Output:

• A cleaned dataset for engagement prediction.

use cuda and load Top 100 social media profiles dataset Remove duplicate user profiles.

OUTPUT:

```
Original number of profiles: 500

DataFrame after removing duplicates based on 'PROFILE':
New number of profiles: 465
```

	N	PROFILE	FOLLOWERS	POSTS	Platform
0	1	Instagram	539446645.0	7202.0	Instagram
1	2	Cristiano Ronaldo	473864939.0	3338.0	Instagram
2	3	Kylie <U+0001F90D>	364542529.0	6935.0	Instagram
3	4	Leo Messi	355790796.0	890.0	Instagram
4	5	Selena Gomez	341579063.0	1828.0	Instagram

Handle missing values

OUTPUT:

```
display(df_social_cleaned.head())
```

	N	PROFILE	FOLLOWERS	POSTS	Platform
0	1	Instagram	539446645.0	7202.0	Instagram
1	2	Cristiano Ronaldo	473864939.0	3338.0	Instagram
2	3	Kylie <U+0001F90D>	364542529.0	6935.0	Instagram
3	4	Leo Messi	355790796.0	890.0	Instagram
4	5	Selena Gomez	341579063.0	1828.0	Instagram

Encode categorical columns

OUTPUT:

```
DataFrame after one-hot encoding 'Platform' column:
```

	N	PROFILE	FOLLOWERS	POSTS	Platform_TikTok	Platform_Twitch	Platform
0	1	Instagram	539446645.0	7202.0	0	0	
1	2	Cristiano Ronaldo	473864939.0	3338.0	0	0	
2	3	Kylie <U+0001F90D>	364542529.0	6935.0	0	0	
3	4	Leo Messi	355790796.0	890.0	0	0	
4	5	Selena Gomez	341579063.0	1828.0	0	0	

```
Shape of DataFrame after encoding:  (465, 8)
```

Task 11 – Weather Dataset Feature Engineering

Task:

Use AI to preprocess a weather dataset for forecasting.

Instructions:

• Convert date column into datetime format.

• Handle missing values in temperature, humidity.

• Create new features (season, week_number).

• Normalize rainfall values.

Dataset link:

https://www.kaggle.com/code/alsamaualalhinai/weather-dataset/input

Expected Output:

• A processed dataset ready for forecasting models.

use cuda and load weather_forecast_data.csv datset and Convert date column into datetime format.

OUTPUT:

```
Available columns: ['Temperature', 'Humidity', 'Wind_Speed', 'Cloud_Cover', 'Pressu
DataFrame loaded using cuDF (on GPU):
```

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	Rain
0	23.720338	89.592641	7.335604	50.501694	1032.378759	rain
1	27.879734	46.489704	5.952484	4.990053	992.614190	no rain
2	25.069084	83.072843	1.371992	14.855784	1007.231620	no rain
3	23.622080	74.367758	7.050551	67.255282	982.632013	rain
4	20.591370	96.858822	4.643921	47.676444	980.825142	no rain

Handle missing values in temperature, humidity.

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	Rain
0	23.720338	89.592641	7.335604	50.501694	1032.378759	rain
1	27.879734	46.489704	5.952484	4.990053	992.614190	no rain
2	25.069084	83.072843	1.371992	14.855784	1007.231620	no rain
3	23.622080	74.367758	7.050551	67.255282	982.632013	rain
4	20.591370	96.858822	4.643921	47.676444	980.825142	no rain

Normalize rainfall values.

```
Rainfall values normalized using Min-Max scaling.
First 5 rows with new 'Rain_normalized' column:
```

	Temperature	Humidity	Wind_Speed	Cloud_Cover	Pressure	Rain	Rain_numerica
0	23.720338	89.592641	7.335604	50.501694	1032.378759	rain	1.
1	27.879734	46.489704	5.952484	4.990053	992.614190	no rain	0.
2	25.069084	83.072843	1.371992	14.855784	1007.231620	no rain	0.
3	23.622080	74.367758	7.050551	67.255282	982.632013	rain	1.
4	20.591370	96.858822	4.643921	47.676444	980.825142	no rain	0.

Task 12 – Movie Ratings Dataset Preparation

Task:

Clean and preprocess a movie ratings dataset using AI-generated scripts.

Instructions:

• Handle missing values in rating, genre.
• Encode categorical variables like genre, language.
• Remove duplicate movie entries.
• Scale rating values between 0 and 1.

Dataset link:
https://www.kaggle.com/datasets/PromptCloudHQ/imdb-data

Expected Output:
• A dataset suitable for recommendation systems.

```
use cuda an d load IMDB-Movie-Data.csv dataset
Clean and preprocess a movie ratings dataset Handle
missing values in rating, genre.
```

OUTPUT:

```
Dataset loaded successfully! Here is the initial info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  ------              --------------  -----
 0   Rank                1000 non-null   int64
 1   Title               1000 non-null   object
 2   Genre               1000 non-null   object
 3   Description         1000 non-null   object
 4   Director            1000 non-null   object
 5   Actors              1000 non-null   object
 6   Year                1000 non-null   int64
 7   Runtime (Minutes)   1000 non-null   int64
 8   Rating              1000 non-null   float64
 9   Votes               1000 non-null   int64
 10  Revenue (Millions)  872 non-null    float64
 11  Metascore           936 non-null    float64
dtypes: float64(3), int64(4), object(5)
memory usage: 93.9+ KB
```

```
Now, let's examine the missing values specifically in the 'Rating' and 'Genre' columns.
```

```
Missing values after handling all numerical columns:
Revenue (Millions)    0
Metascore             0
dtype: int64

Dataset info after handling all missing numerical values:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  ------              --------------  -----
 0   Rank                1000 non-null   int64
 1   Title               1000 non-null   object
 2   Genre               1000 non-null   object
 3   Description         1000 non-null   object
 4   Director            1000 non-null   object
 5   Actors              1000 non-null   object
 6   Year                1000 non-null   int64
 7   Runtime (Minutes)   1000 non-null   int64
 8   Rating              1000 non-null   float64
 9   Votes               1000 non-null   int64
 10  Revenue (Millions)  1000 non-null   float64
 11  Metascore           1000 non-null   float64
dtypes: float64(3), int64(4), object(5)
memory usage: 93.9+ KB
/tmp/ipykernel_8404/3991703251.py:2: FutureWarning: A value is trying to be set on
The behavior will change in pandas 3.0. This inplace method will never work because

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method
```

```
Encode categorical variables like genre, language.
```

OUTPUT:

DataFrame after one-hot encoding genres:

	Rank	Title	Description	Director	Actors	Year	Runtime (Minutes)	Ra
0	1	Guardians of the Galaxy	A group of intergalactic criminals are forced ...	James Gunn	Chris Pratt, Vin Diesel, Bradley Cooper, Zoe S...	2014	121	
1	2	Prometheus	Following clues to the origin of mankind, a te...	Ridley Scott	Noomi Rapace, Logan Marshall-Green, Michael Fa...	2012	124	
2	3	Split	Three girls are kidnapped by a man with a diag...	M. Night Shyamalan	James McAvoy, Anya Taylor-Joy, Haley Lu Richar...	2016	117	
3	4	Sing	In a city of humanoid animals, a hustling thea...	Christophe Lourdelet	Matthew McConaughey,Reese Witherspoon, Seth Ma...	2016	108	
4	5	Suicide Squad	A secret government agency recruits some of th...	David Ayer	Will Smith, Jared Leto, Margot Robbie, Viola D...	2016	123	

Remove duplicate movie entries.

OUTPUT:

Number of duplicate movie entries found (based on Title and Year): 0
DataFrame shape after removing duplicates: (1000, 31)

First 5 rows of the DataFrame after removing duplicates:

	Rank	Title	Description	Director	Actors	Year	Runtime (Minutes)	Ra
0	1	Guardians of the Galaxy	A group of intergalactic criminals are forced ...	James Gunn	Chris Pratt, Vin Diesel, Bradley Cooper, Zoe S...	2014	121	
1	2	Prometheus	Following clues to the origin of mankind, a te...	Ridley Scott	Noomi Rapace, Logan Marshall-Green, Michael Fa...	2012	124	
2	3	Split	Three girls are kidnapped by a man with a diag...	M. Night Shyamalan	James McAvoy, Anya Taylor-Joy, Haley Lu Richar...	2016	117	
3	4	Sing	In a city of humanoid animals, a hustling thea...	Christophe Lourdelet	Matthew McConaughey,Reese Witherspoon, Seth Ma...	2016	108	
4	5	Suicide	A secret government agency	David Ayer	Will Smith, Jared Leto, Margot Robbie	2016	123	

Scale rating values between 0 and 1.

OUTPUT:

First 5 rows with the new scaled Rating column:

	Rating	Rating_Scaled
0	8.1	0.873239
1	7.0	0.718310
2	7.3	0.760563
3	7.2	0.746479
4	6.2	0.605634

Original Rating min: 1.90, max: 9.00
Scaled Rating min: 0.00, max: 1.00